

▼ Imports & private packages

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

▶ Private packages via Google Drive

[] ↳ 4 cells hidden

▶ Imports

[] ↳ 1 cell hidden

▼ Training routine

```
1 def epoch(net, train_loader, loss_func, optimizer, batch_size, seq_len, measurement='G'
2     epoch_losses = np.zeros([len(train_loader), ])
3     correct = 0
4     incorrect = 0
5
6     for i, d in enumerate(train_loader):
7         # Scale inputs and get predictions
8         inp = d[:, :, :-1].squeeze(1)
9         inp[:, :, 0] /= 1000
10
11         # Add random noise to heat capacity values to simulate real measurements
12         #d[:, :, 1] += torch.normal(mean=0, std=1.25, size=(d.shape[0], d.shape[1], ))
13
14         predictions = net(inp.float().to(device))
15
16         # Each batch consists of measurement batches, where seq_len measurements are pu
17         # measurement batch, every measurement has the same label as it needs to be fro
18         # This leads to a target array where the target class is contained for each mea
19         # batch. With this, CrossEntropyLoss would not work as the predictions are made
20         # batch and CEL therefore expects only on class label per measurement batch. Th
21         # element of the last dimension of d is considered as target (all the entries i
22         # same anyways so it could be any entry)
23         targets = d[:, :, -1][:, 0].long().to(device)
24
25         correct += (predictions.argmax(dim=-1) == targets).sum().item()
26         incorrect += len(targets) - (predictions.argmax(dim=-1) == targets).sum().item(
27
```

```
28     # Calculate the loss
29     loss = loss_func(predictions, targets)
30     epoch_losses[i] = loss
31
32     # Backward step
33     net.zero_grad()
34     loss.backward()
35     optimizer.step()
36
37     accuracy = correct/(correct + incorrect)
38     return epoch_losses.mean(), accuracy

1 from torch.nn.modules.activation import Softmax
2 # Training and testing routines
3
4 def train(net, optimizer, train_loader, test_loader, batch_size, seq_len, measurement='
5     loss_func = nn.CrossEntropyLoss()
6
7     losses = np.zeros([nr_epochs, ])
8     accuracies = np.zeros([nr_epochs, ])
9
10    best_loss, _ = epoch(net, train_loader, loss_func, optimizer, batch_size, seq_len,
11    best_net = net
12
13    for i in range(nr_epochs):
14        losses[i], accuracies[i] = epoch(net, train_loader, loss_func, optimizer, batch
15        if losses[i] < best_loss:
16            best_net = net
17            best_loss = losses[i]
18
19        if i % 10 == 0:
20            print('Epoch ', i)
21            print('Loss: ', losses[i])
22            print('Training accuracy: ', accuracies[i])
23            #test(best_net, test_loader)
24
25    return best_net, losses, accuracies
26
27 def test(net, test_loader):
28     correct = 0
29     incorrect = 0
30
31     for d in test_loader:
32         inp = d[:, :, :-1]#.squeeze(1)
33         inp[:, :, 0] /= 1000
34         predictions = net(inp.float().to(device)).squeeze()#.argmax()
35         targets = d[:, :, -1][:, 0].long().to(device)
36
37         correct += (predictions.argmax(dim=-1) == targets).sum().item()
38         incorrect += len(targets) - (predictions.argmax(dim=-1) == targets).sum().item(
39         #print(predictions.argmax(dim=-1), targets, Softmax()(predictions).amax(dim=-1)
40
41     accuracy = correct/(correct + incorrect)
42     print('Test accuracy: ', accuracy)
```

```
43
44     return accuracy
```

▼ Run

▼ Network definition

```
1 # Hyperparameters
2 seq_len = 5
3 measurement = 'C'
4 batch_size = 256
5 nr_epochs = 500
6 lr = 0.001
7 hidden_size_linear = 128
8 hidden_layers = 2
9 step = 0.05

1 def create_data_loaders(elements, splits, seq_len, validation=False, measurement='G', u
2     dc = DatasetCreator(elements=elements, splits=splits, validation=validation, seq_l
3     train_dataset, test_dataset, val_dataset = dc.get_datasets()
4
5     # Create the DataLoaders
6     train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=shuffle)
7     test_loader = DataLoader(test_dataset, batch_size=1, shuffle=shuffle)
8     if val_dataset:
9         val_loader = DataLoader(val_dataset)
10    else:
11        val_loader = None
12
13    return train_loader, test_loader, val_loader

1 # Create the network
2 t = True
3 net = ElementClassifier(train, in_features=seq_len * 2, hidden_size_linear=hidden_size_
4
5 # Check if cuda is available and send net to device
6 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
7
8 net = net.to(device)
9
10 optimizer = Adam(net.parameters(), lr=lr)
```

▼ Train

```
1 elements = None
2
```

```
3 train_loader, test_loader, val_loader = create_data_loaders(elements, (0.8, 0.2), seq_1
```

```
    Dataset shape: (411319, 5, 3)
```

```
1 best_net, losses, accuracies = train(net, optimizer, train_loader, test_loader, batch_s
```

```
2 test(best_net, test_loader)
```

```
    Loss: 0.07865436186036877
```

```
    Training accuracy: 0.9723426343057335
```

```
    Epoch 320
```

```
    Loss: 0.0778072924285376
```

```
    Training accuracy: 0.9730890136366178
```

```
    Epoch 330
```

```
    Loss: 0.07234535236740351
```

```
    Training accuracy: 0.9745623226741288
```

```
    Epoch 340
```

```
    Loss: 0.07166698922809704
```

```
    Training accuracy: 0.9751968666655321
```

```
    Epoch 350
```

```
    Loss: 0.07329059858635756
```

```
    Training accuracy: 0.9744723681619376
```

```
    Epoch 360
```

```
    Loss: 0.06884146940096976
```

```
    Training accuracy: 0.9762349903602799
```

```
    Epoch 370
```

```
    Loss: 0.06597228690198904
```

```
    Training accuracy: 0.9770713242033555
```

```
    Epoch 380
```

```
    Loss: 0.06901550271670615
```

```
    Training accuracy: 0.9760088884782857
```

```
    Epoch 390
```

```
    Loss: 0.07293795930491453
```

```
    Training accuracy: 0.9750777377169545
```

```
    Epoch 400
```

```
    Loss: 0.06974695115989131
```

```
    Training accuracy: 0.9758605850933217
```

```
    Epoch 410
```

```
    Loss: 0.06736178178277606
```

```
    Training accuracy: 0.9765850835969163
```

```
    Epoch 420
```

```
    Loss: 0.07149237042446387
```

```
    Training accuracy: 0.9757001256931968
```

```
    Epoch 430
```

```
    Loss: 0.07081095838838002
```

```
    Training accuracy: 0.9763881561513084
```

```
    Epoch 440
```

```
    Loss: 0.06990428069043834
```

```
    Training accuracy: 0.9762714584057629
```

```
    Epoch 450
```

```
    Loss: 0.06859079741683917
```

```
    Training accuracy: 0.9764659546483386
```

```
    Epoch 460
```

```
    Loss: 0.0670052903234772
```

```
    Training accuracy: 0.9772609580398669
```

```
    Epoch 470
```

```
    Loss: 0.0672400914410485
```

```
    Training accuracy: 0.9771199482639995
```

```
    Epoch 480
```

```
    Loss: 0.06531707772162297
```

```
    Training accuracy: 0.9777423362402418
```

```
    Epoch 490
```

Epoch: 490

Loss: 0.06907681258250864

Training accuracy: 0.9770664617972912

Test accuracy: 0.9806318694667691

0.9806318694667691

```
1 test(best_net, test_loader)
```

```
1 torch.save(best_net, 'ElementClassifier_9806.pth')
```

```
2 # PLOTS!!!
```

```
1 print(type(losses))
```

```
2 plt.plot(losses)
```

```
3 plt.xlabel('Epochs'), plt.ylabel('Loss')
```

```
4 plt.title('Losses over time')
```

```
5 plt.grid()
```

```
6 plt.show()
```

```
7 plt.plot(accuracies)
```

```
8 plt.xlabel('Epochs'), plt.ylabel('Training accuracy')
```

```
9 plt.title('Training accuracy over time')
```

```
10 plt.grid()
```

```
11 plt.show()
```

▼ Test on Barin data

```
1 net = torch.load('/content/ElementClassifier_9782_3.pth').to(device)
```

```
1 inp = torch.tensor([[ .3,  11.403],  
2                      [ .7,  22.25],  
3                      [ .8,  23.364],  
4                      [ .9,  24.248],  
5                      [1.0,  24.979]]).to(device)
```

```
1 out = net(inp)
```

```
1 print(Encoder()(out.argmax(dim=-1).item()))  
2 print(Softmax(dim=-1)(out).amax(dim=-1))
```

```
B  
tensor([1.0000], grad_fn=<AmaxBackward0>)
```

```
1 test(net, test_loader)
```

```
Test accuracy: 0.9772111058201572  
0.9772111058201572
```

✓ 0s completed at 7:20 PM

